Creating a High-Availability Web Application Infrastructure on AWS with VPC, Auto Scaling, and Application Load Balancer

Stephanie Tabares · Follow

6 min read · Feb 21, 2023

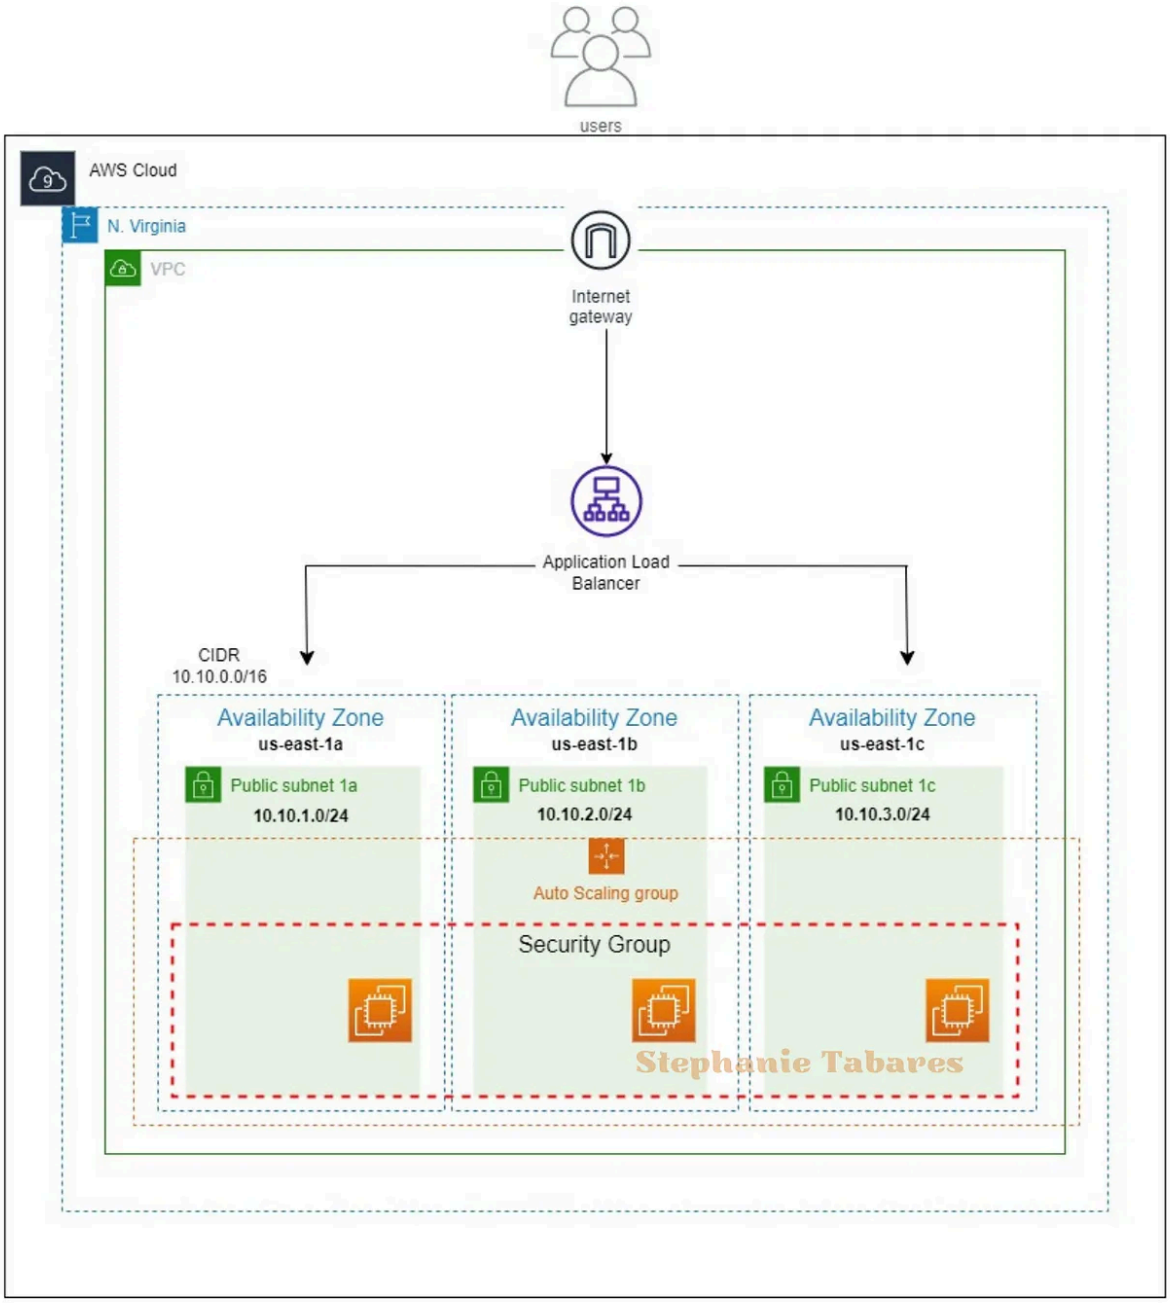

Welcome to project #5! In today's digital age, having a reliable and scalable web application infrastructure is crucial for any business. In this project, we will be creating a Virtual Private Cloud (VPC) on Amazon Web Services (AWS) with three public subnets and an autoscaling group to ensure high availability and scalability of our web application. We'll also be configuring an Application Load Balancer to distribute traffic to the autoscaling group, and implementing security groups to ensure the safety and security of our infrastructure. So, let's get started and build an infrastructure that can handle any amount of traffic with ease!

FOUNDATIONAL

1. For this project you MUST create a diagram of your AWS architecture and use it as your Preview Image for your Medium.
2. Create a VPC with cidr 10.10.0.0/16
3. Create three public subnets with 10.10.1.0/24 & 10.10.2.0/24 & 10.10.3.0/24
4. Create an autoscaling group using t2.micro instances. All instances should have apache installed on each instance with the ability to check any random IP address and be able to produce a test page. Ensure the autoscaling group is using the public subnets from #2.
5. The autoscaling min and max should be 2 and 5.
6. Create an Application Load Balancer to distribute traffic to the autoscaling group.
7. Create web server security group that allows inbound traffic from HTTP from your Application Load Balancer.
8. Create a load balancer security group that allows inbound traffic from HTTP from 0.0.0.0/0.

Creating a VPC

1. In the search bar type VPC -> Create VPC -> VPC and more -> Name your VPC -> Add CIDR block 10.10.0.0/16

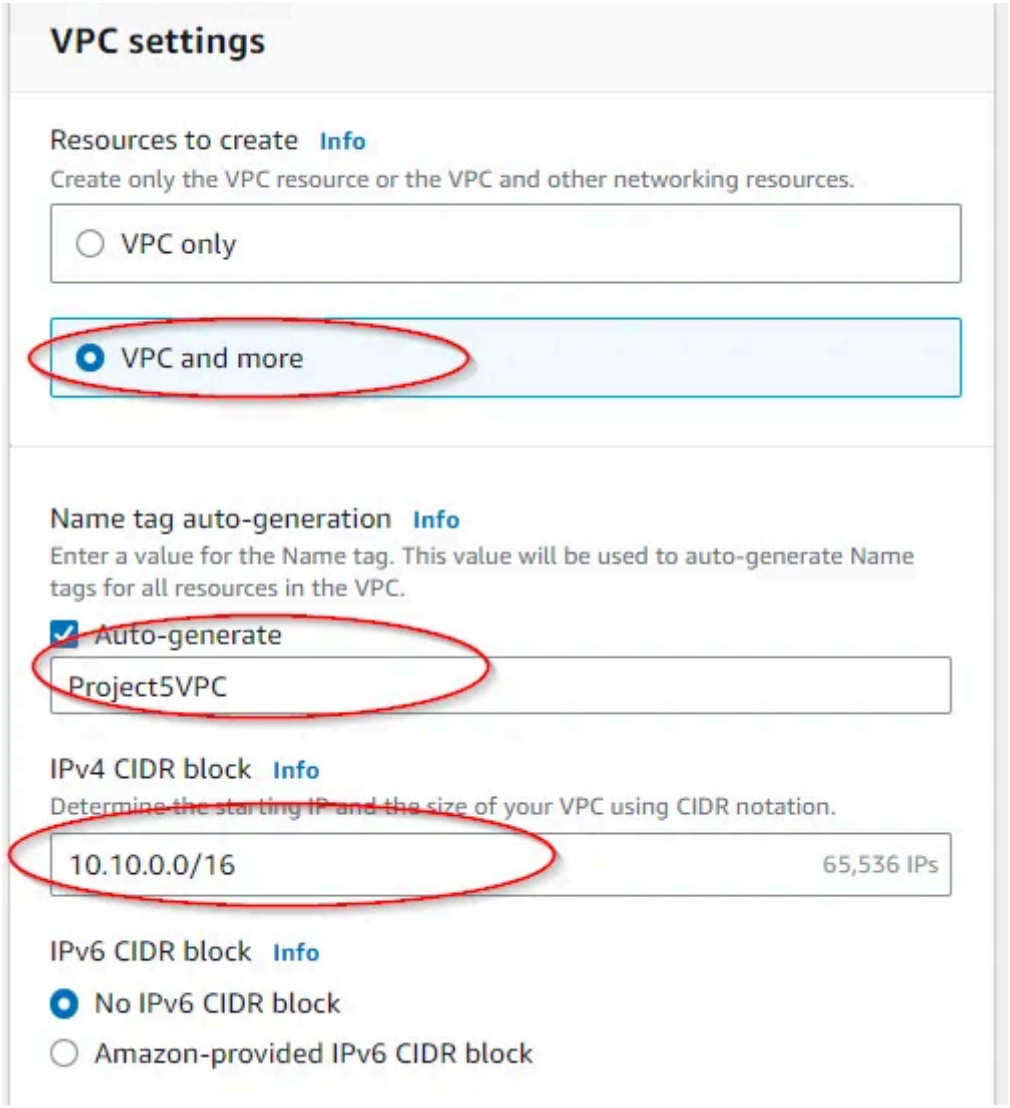

2. Deploying subnets in different availability zones provides high availability and fault tolerance for your applications. This is because if one availability zone becomes unavailable, the other two can continue to handle traffic and requests, preventing any disruption to your application or service.

-Create three public subnets with 10.10.1.0/24 & 10.10.2.0/24 & 10.10.3.0/24

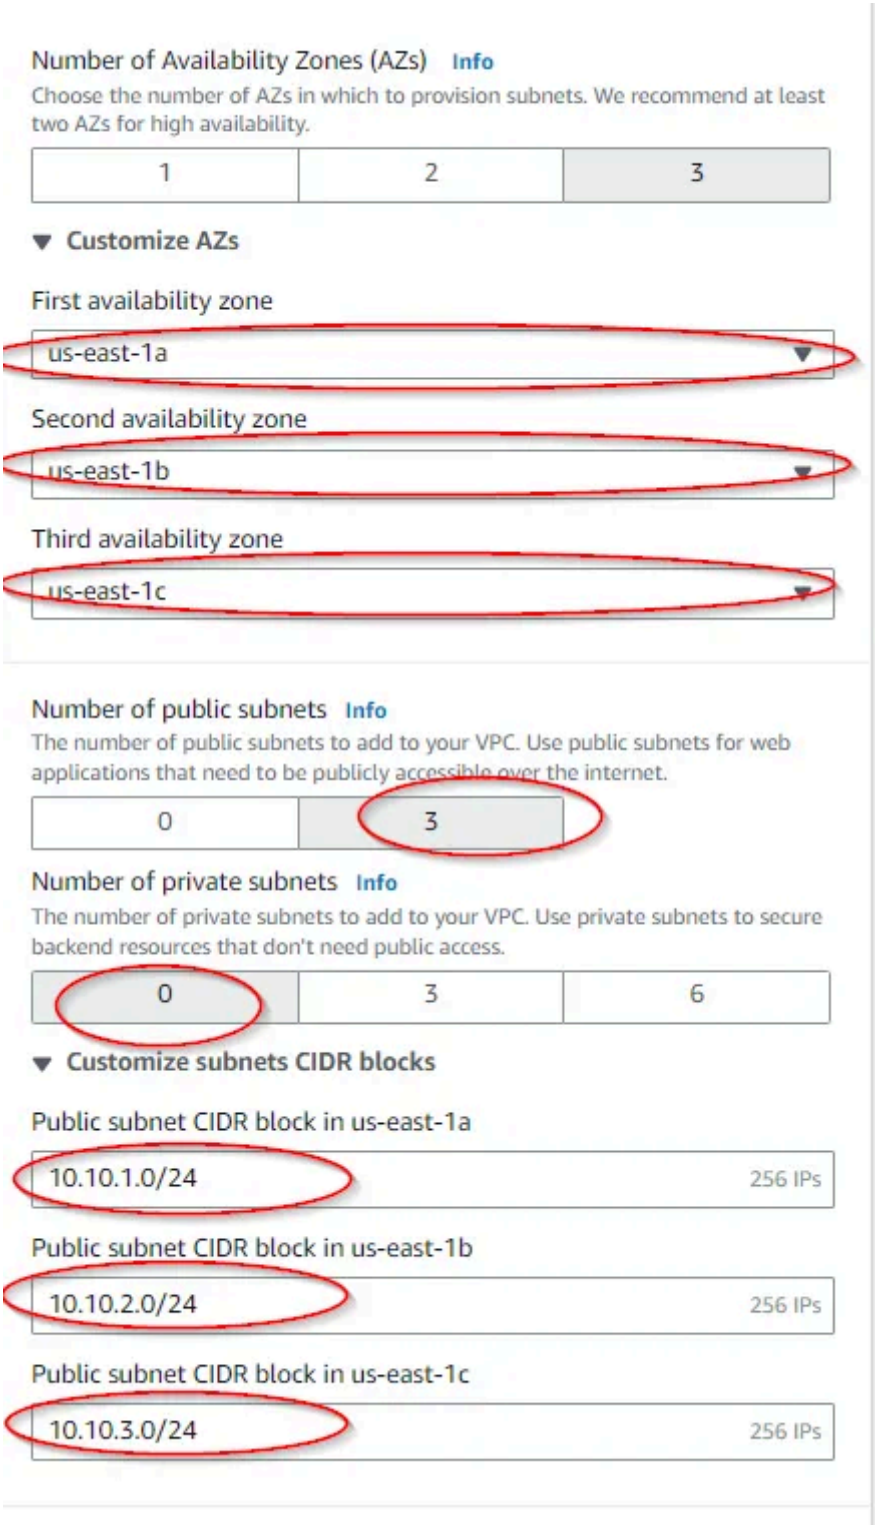
Number of Availability Zones (AZs) Info
Choose the number of AZs in which to provision subnets. We recommend at least two AZs for high availability.
1
2
3
Customize AZs
First availability zone
us-east-1a
Second availability zone
us-east-1b
Third availability zone
us-east-1c
Number of public subnets Info
The number of public subnets to add to your VPC. Use public subnets for web applications that need to be publicly accessible over the internet.
0
3
Number of private subnets Info
The number of private subnets to add to your VPC. Use private subnets to secure backend resources that don't need public access.
0
3
6
Customize subnets CIDR blocks
Public subnet CIDR block in us-east-1a
10.10.1.0/24
256 IPs
Public subnet CIDR block in us-east-1b
10.10.2.0/24
256 IPs
Public subnet CIDR block in us-east-1c
10.10.3.0/24
256 IPs

3. Create VPC

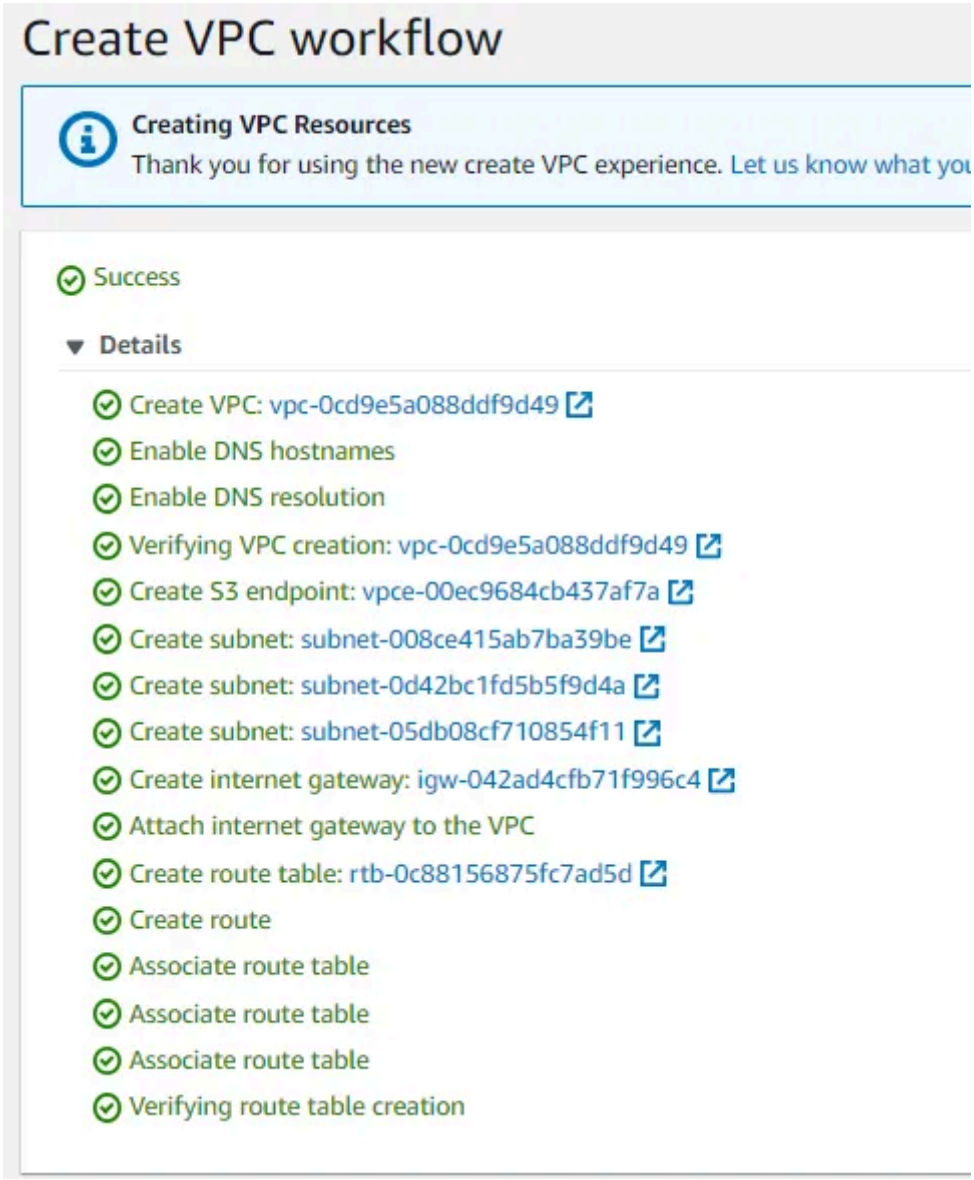

Congrats! Now let's move on to the next step!

Create a launch template

1. In the search bar type EC2 -> Scroll down to Instances -> Launch templates -> Select Create launch template
2. Name launch template. For AMI select Amazon Linux. For instance type select t2.micro. Select your keypair name.

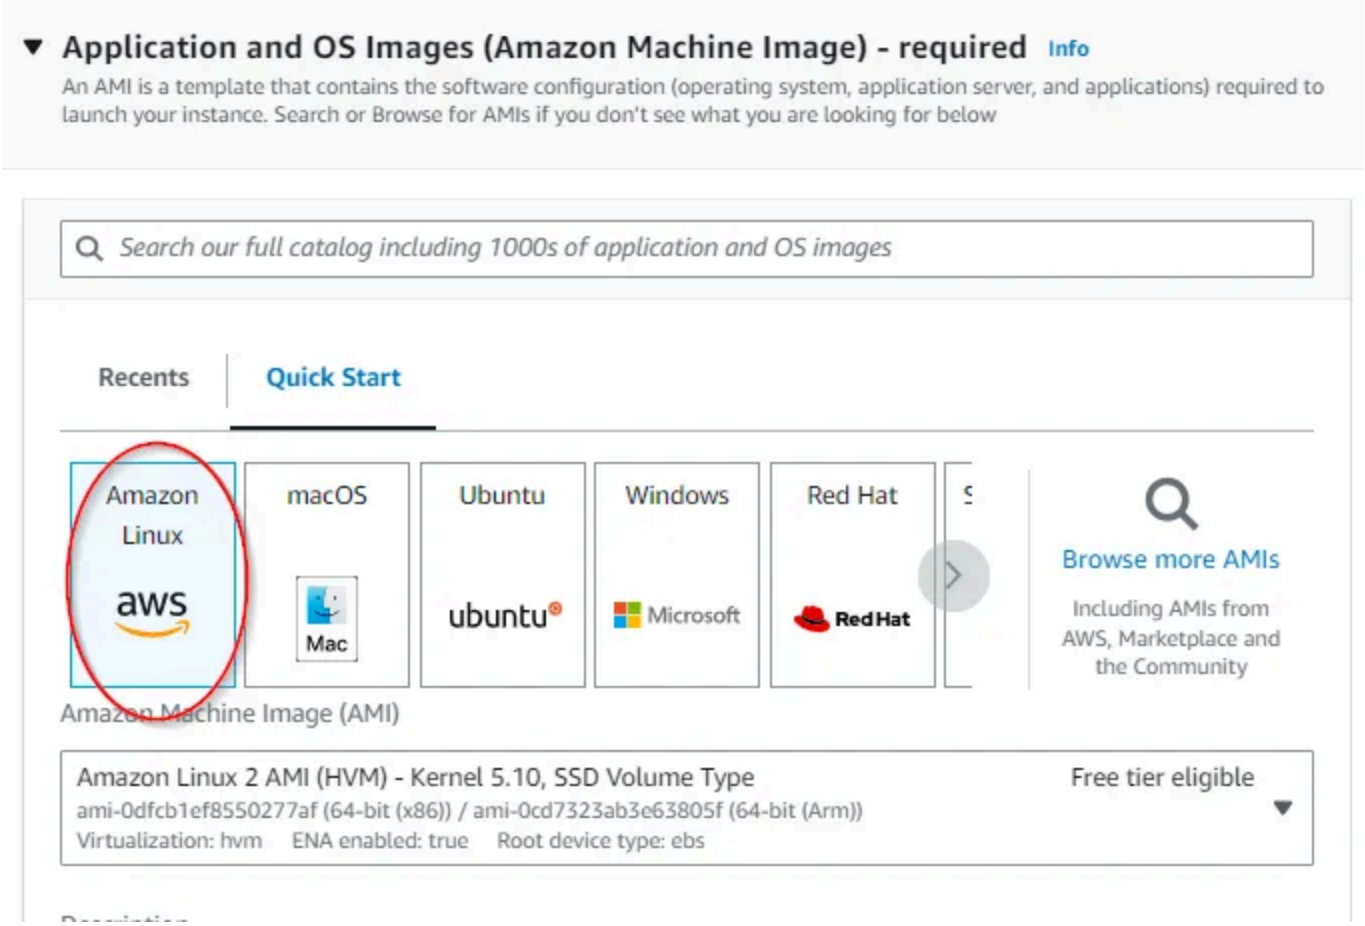

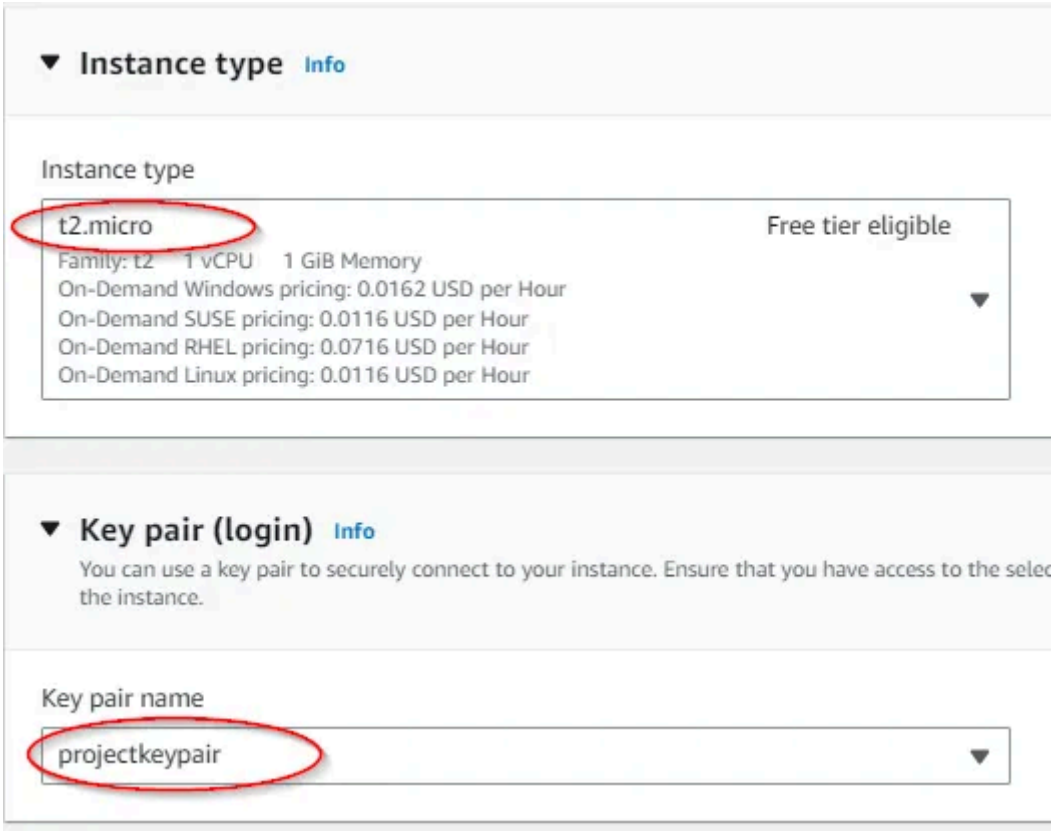

3. Under Network settings select Create security group -> Name your security group -> Allow SSH and HTTP -> Select your VPC (it's

automatically in default so you have to change it)

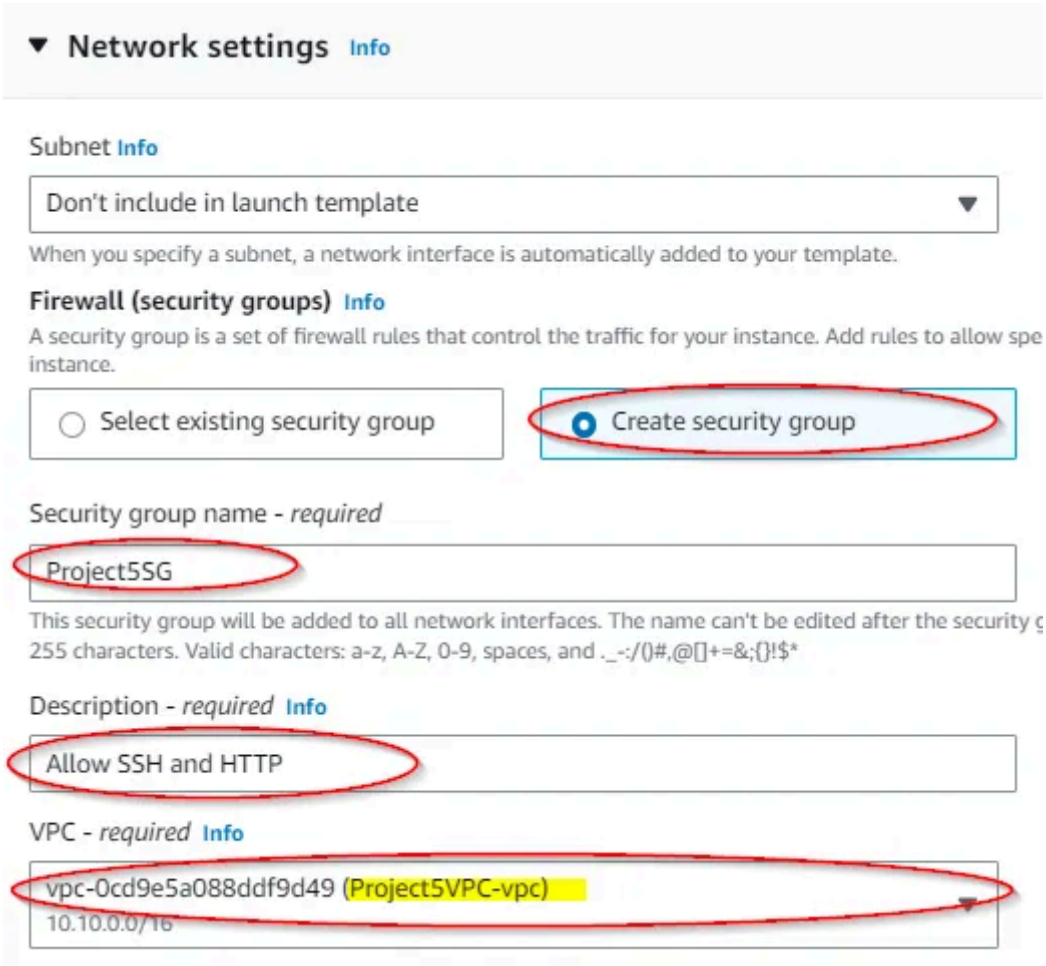

4. For inbound security rules add SSH and HTTP

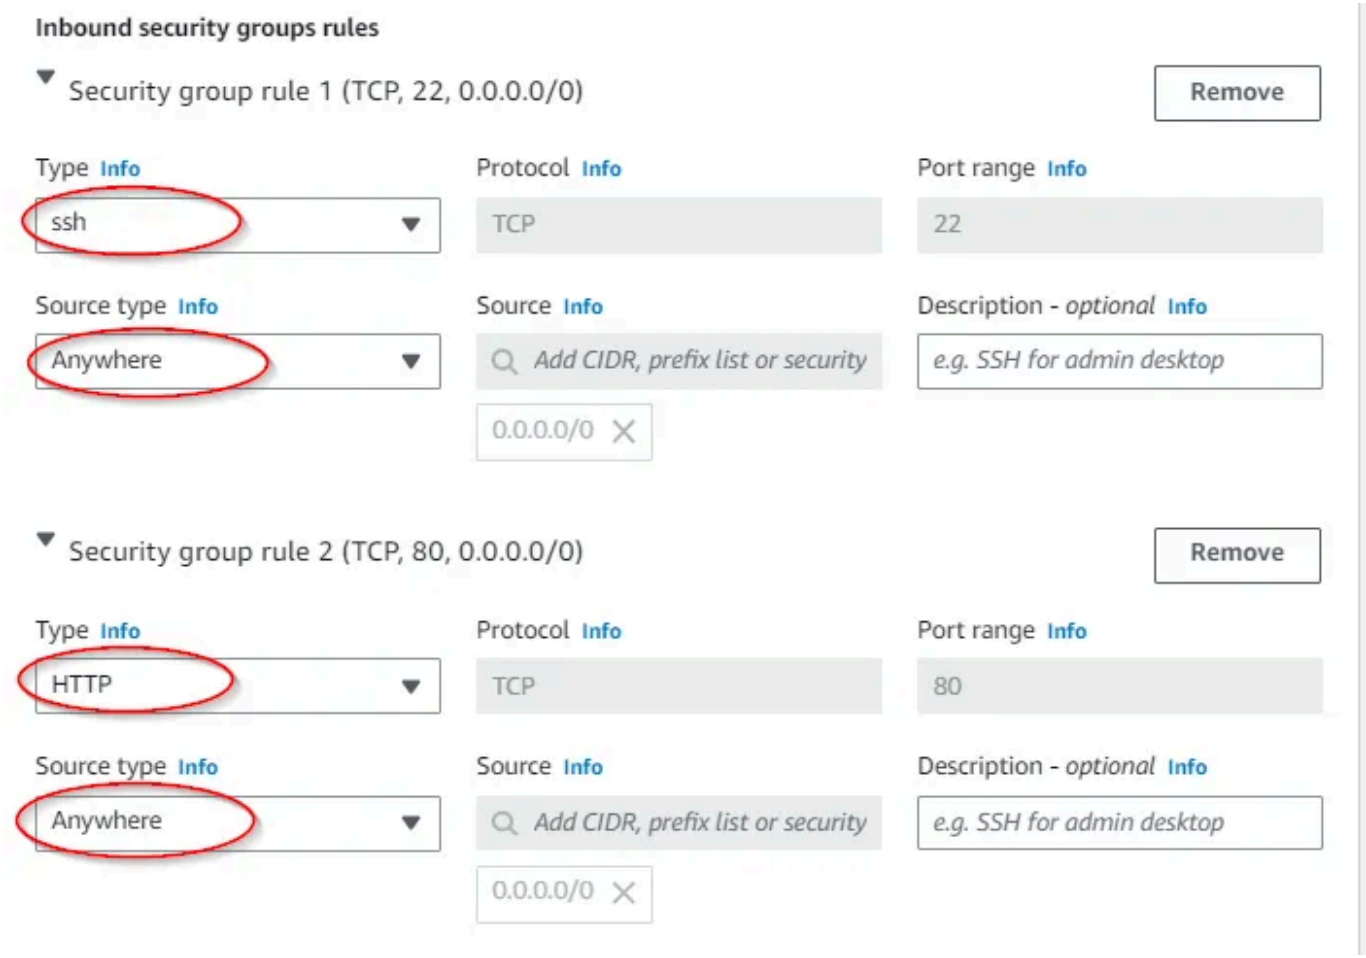

5. For Advanced network configuration select Enable for Auto-assign public IP. A public IP address is necessary for instances in a public subnet to communicate with the internet, receive incoming traffic, and respond to requests from external clients or users.

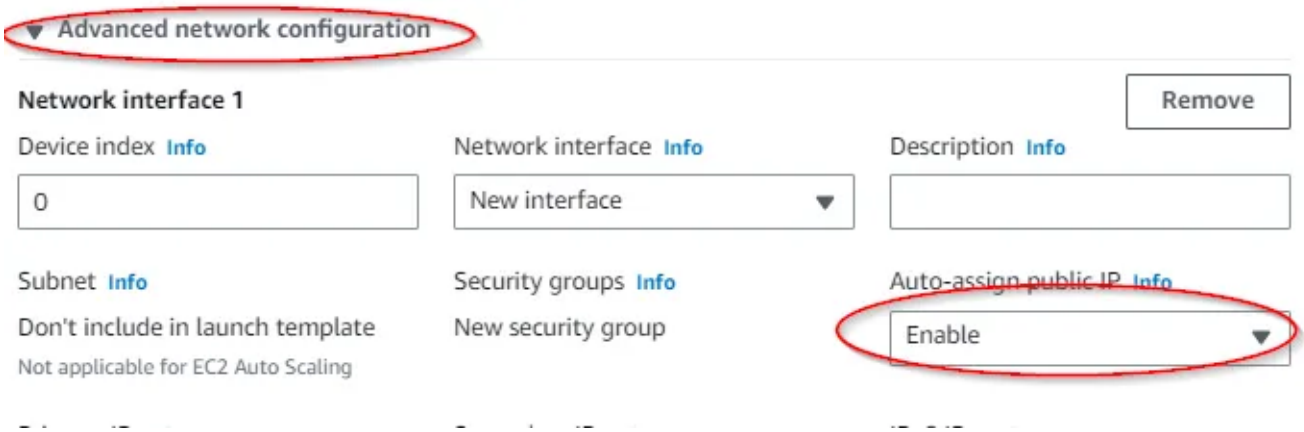

6. In Advanced details scroll to the bottom until you see User data and paste following the command

```
#!/bin/bash
yum update -y
yum install -y httpd
systemctl start httpd
systemctl enable httpd
EC2AZ=$(curl -s http://169.254.169.254/latest/meta-data/placement/availabilit
echo '<center><h1>This Amazon EC2 instance is located in Availability Zone: A
sed "s/AZID/$EC2AZ/" /var/www/html/index.txt >
/var/www/html/index.html
```

7. Create Launch template

Create an Auto Scaling Group

1. Type EC2 in the search bar -> On the left-hand side locate Auto Scaling -> Auto Scaling Groups -> Select Create Auto Scaling Group -> Name group -> Select launch template we just created

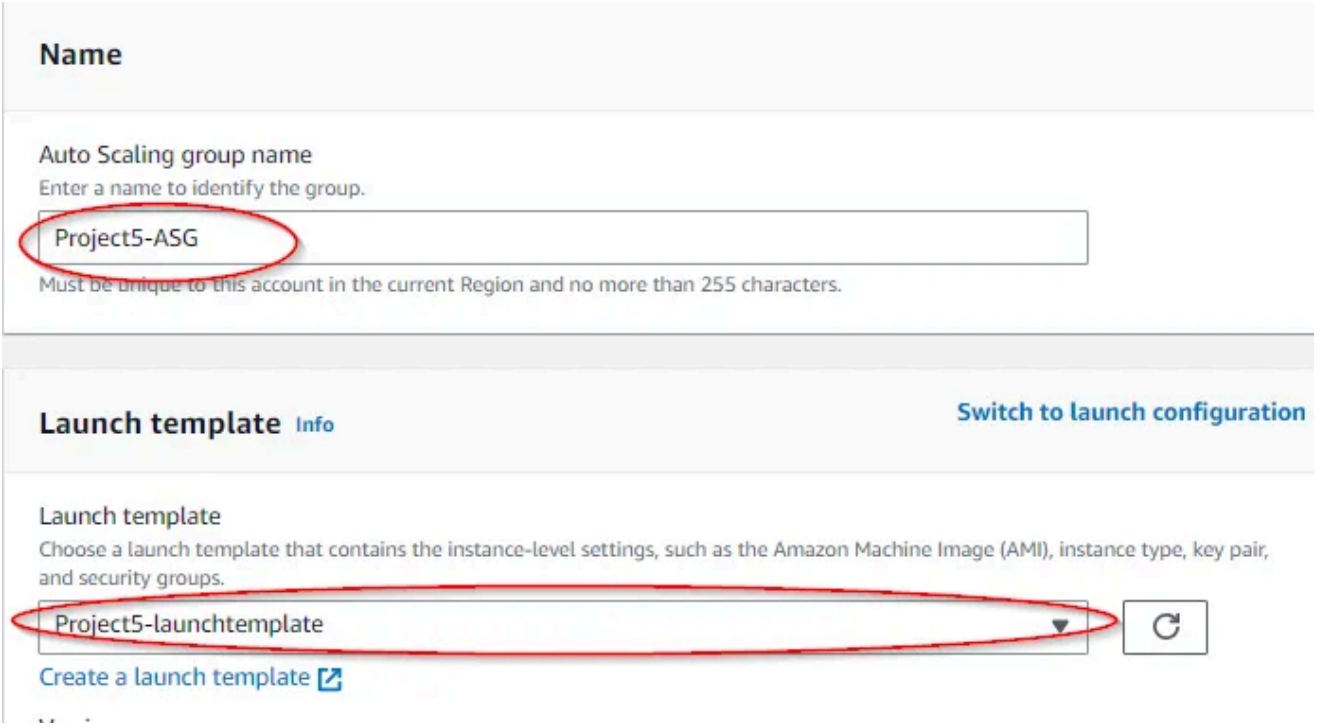

2. Select the VPC we created earlier -> Select all Availability Zones and subnets -> Select next

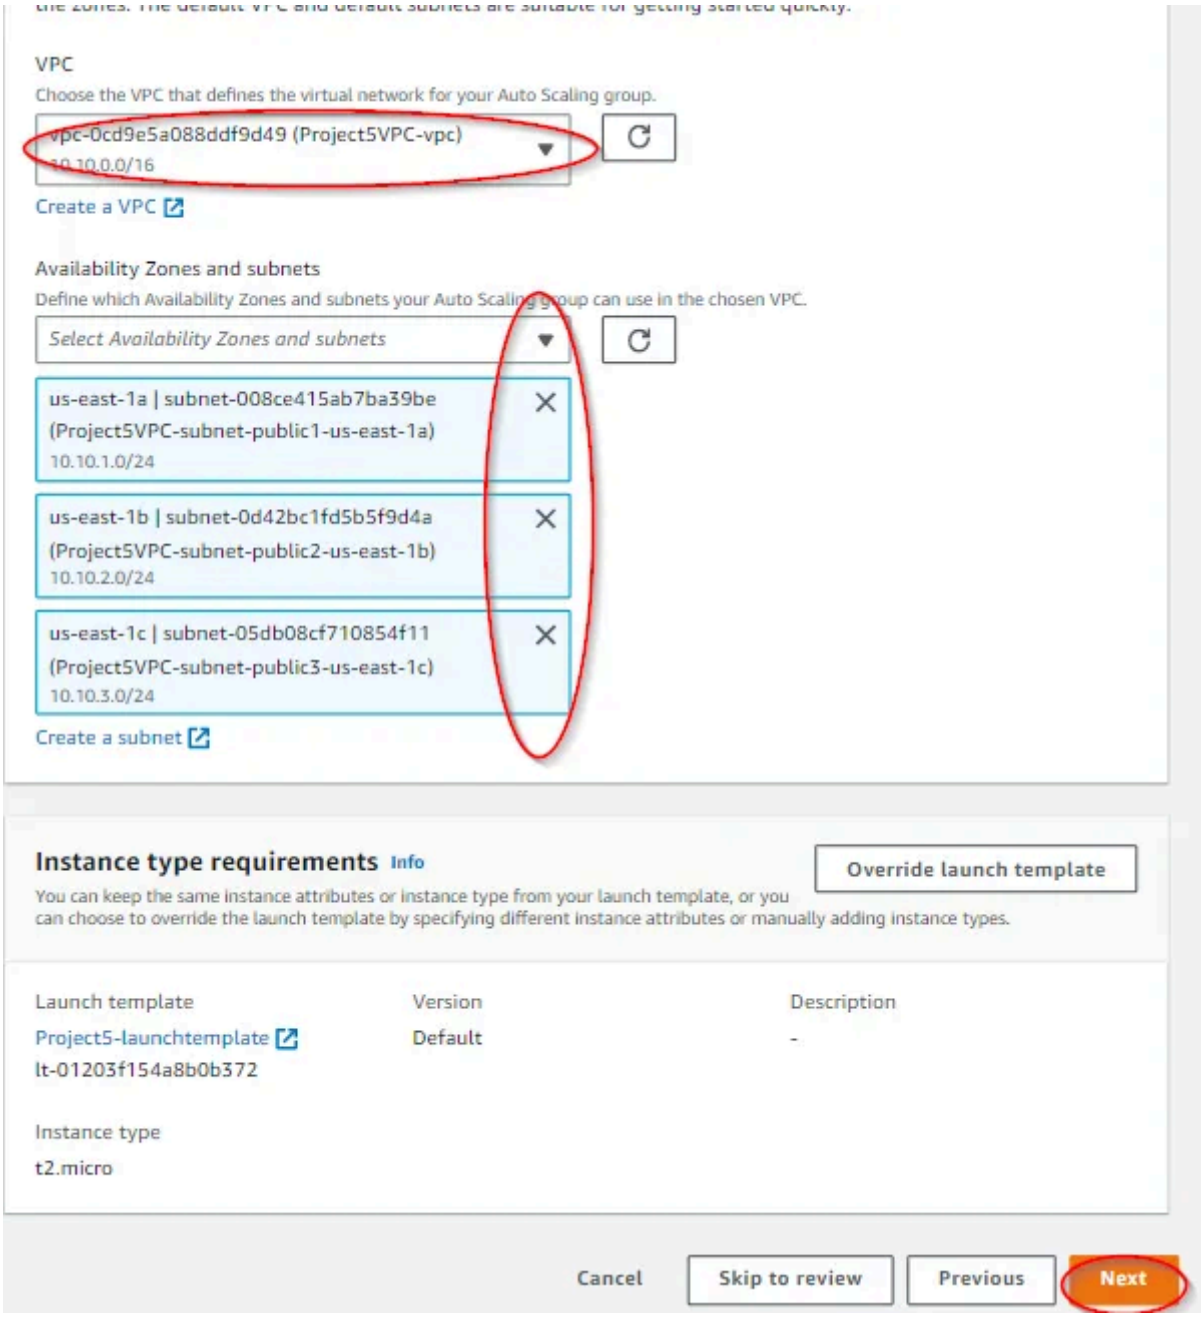

3. We attach a new load balancer in order to distribute incoming traffic evenly across all the instances in the group, ensuring that no single instance becomes overwhelmed or overloaded with requests. An internet-facing load balancer has a public IP address, which clients on the internet can use to connect to your application.

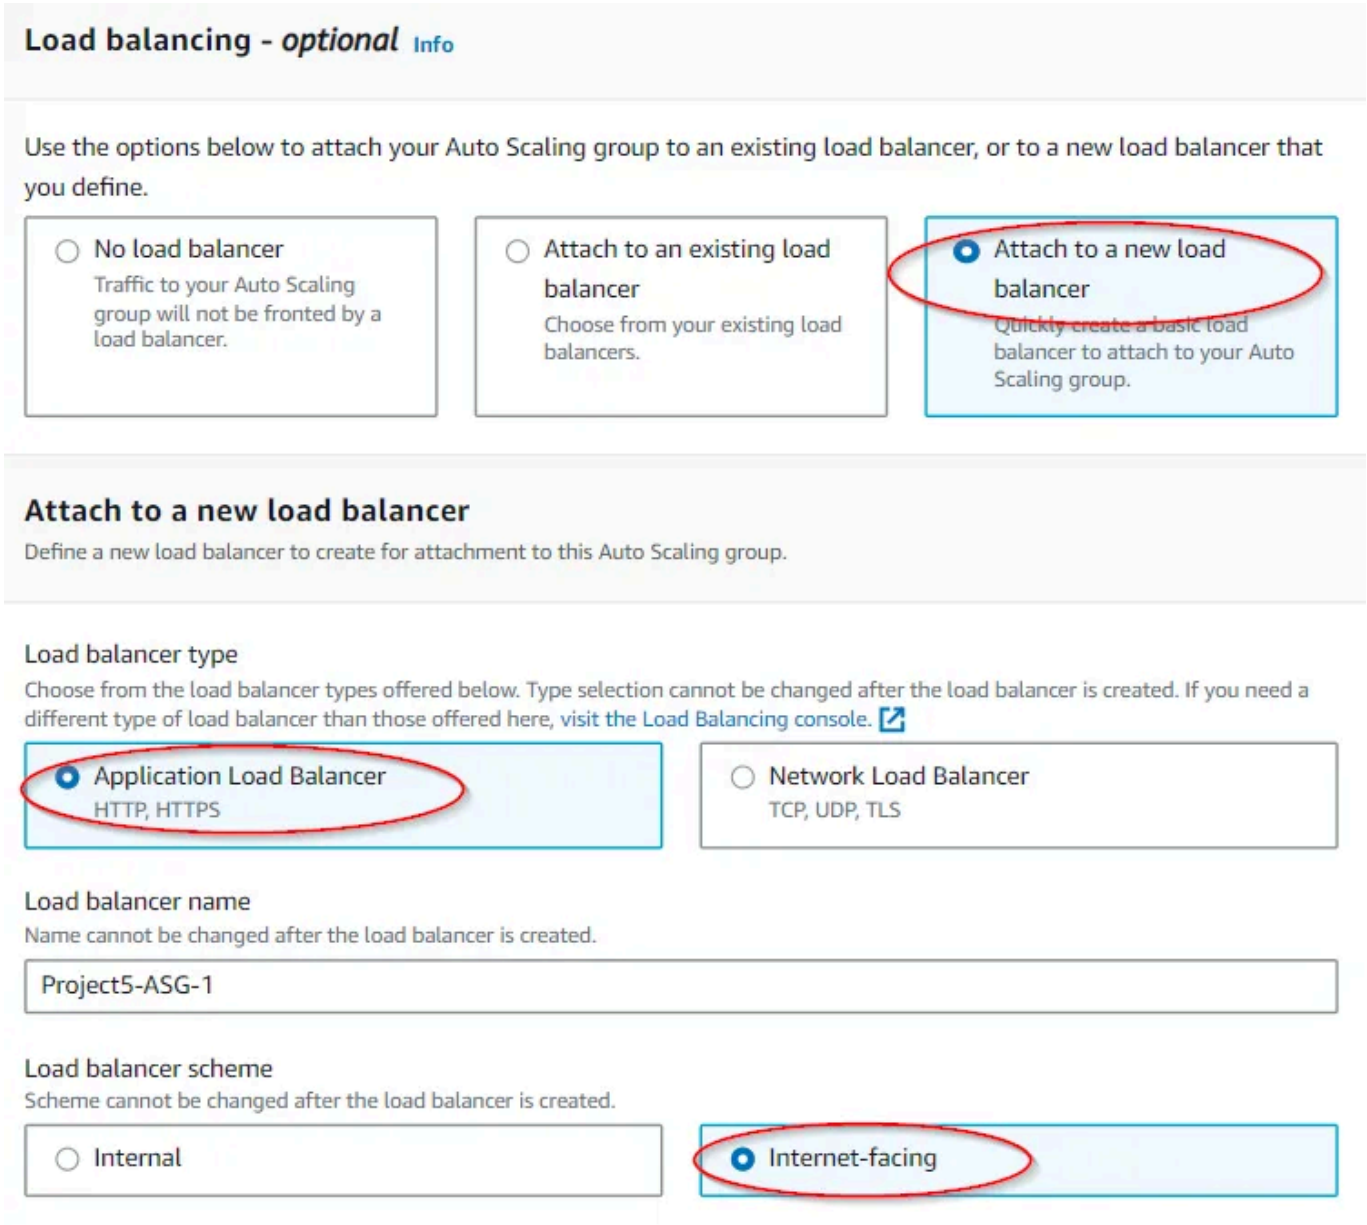

4. When you create a listener and routing rule on your load balancer, you specify which target group(s) should receive the incoming traffic.

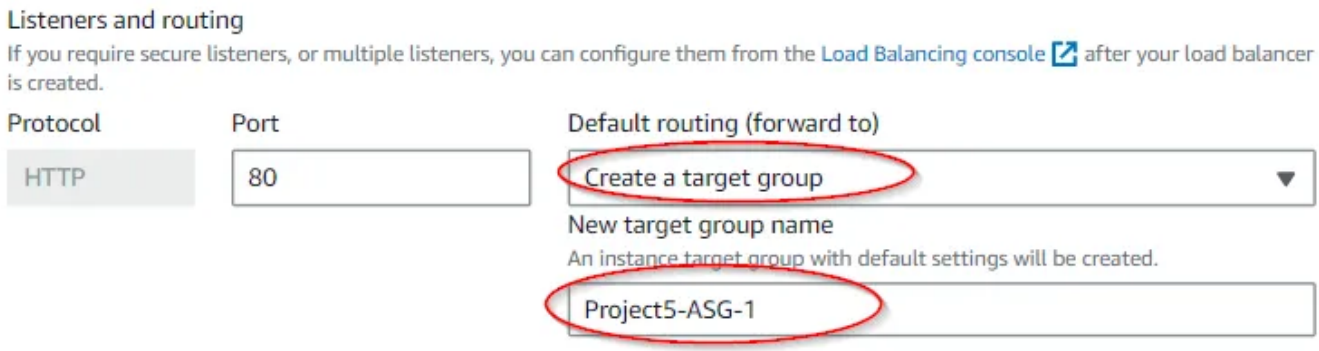

5. Enabling group metrics collection with CloudWatch for your load balancer allows you to monitor the performance of your load balancer and its associated resources.

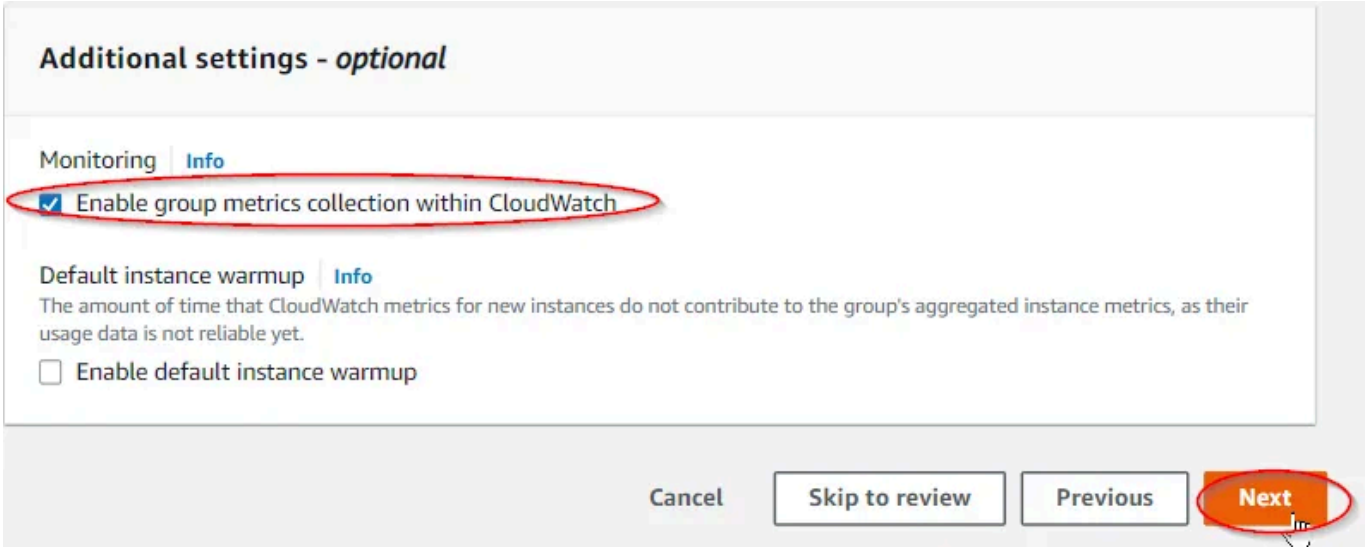

6. The Auto Scaling minimum should be 2 and the maximum 5.

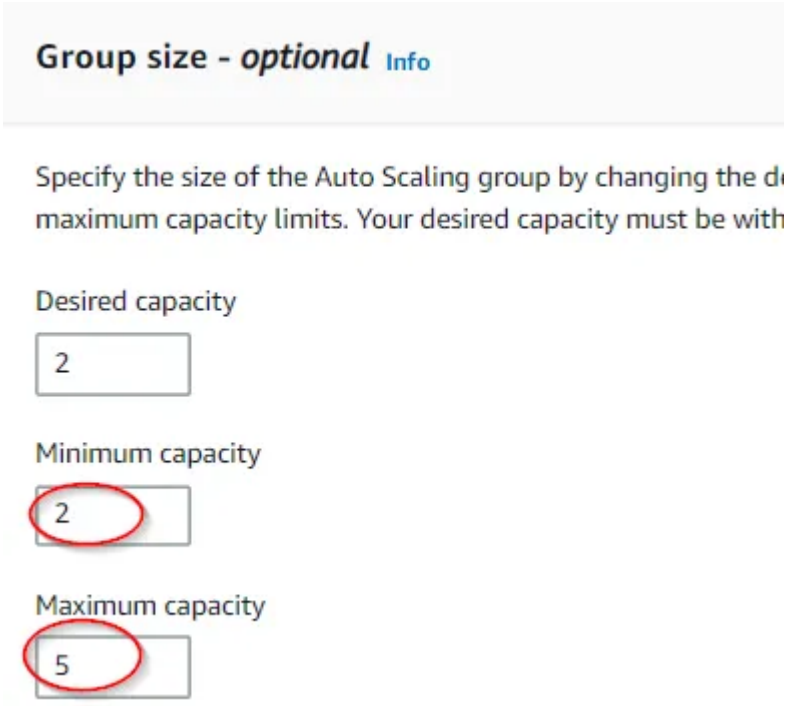

7. Click Next until you locate Create Auto Scaling Group

Let’s see if our Instances are up and running!

1. In the search bar type EC2 -> Instances

2. Locate the Public IPv4 address -> Open browser -> http://34.201.53.174

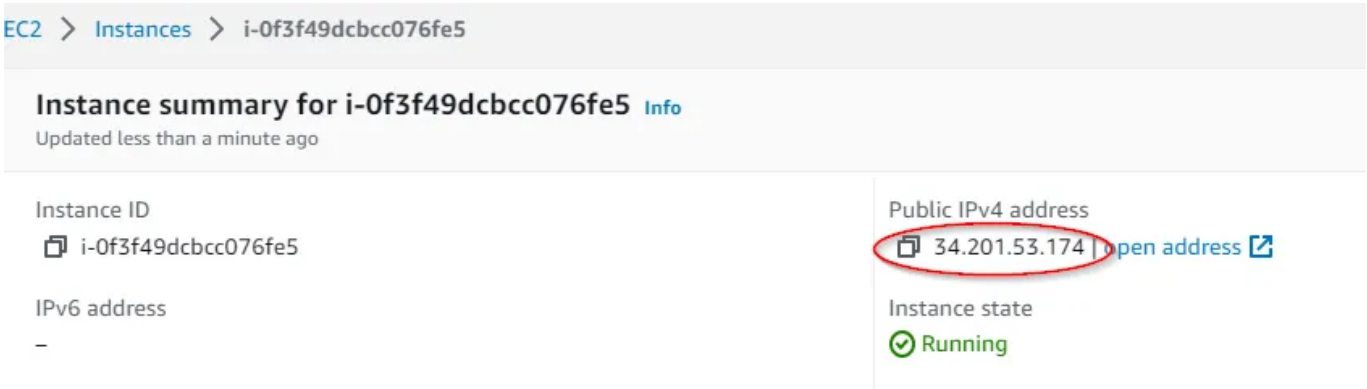

3. Everything looks good to go!

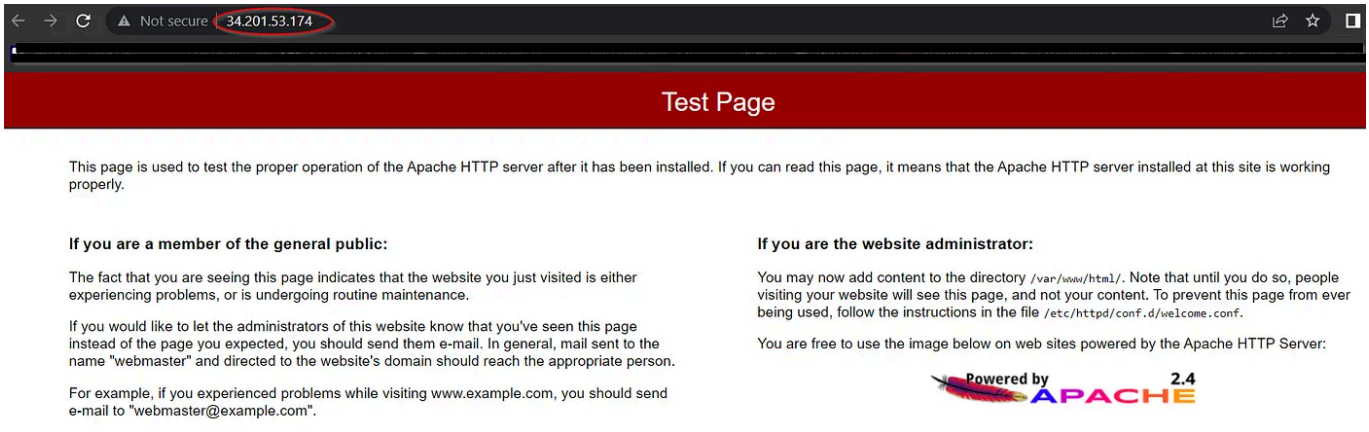

Advanced:

Add a target policy for the ASG to scale after cpu utilization is above 50%. After the autoscaling group has been created, find a stress tool to be able to stress an instance above 50% to see if your scaling policy works! After the autoscaling group has been created, find a stress tool to be able to stress an instance above 50% to see if your scaling policy works!

1. In the search bar type EC2-> Scroll down to Auto Scaling groups > Select group > Go to Automatic scaling > Create dynamic scaling policy

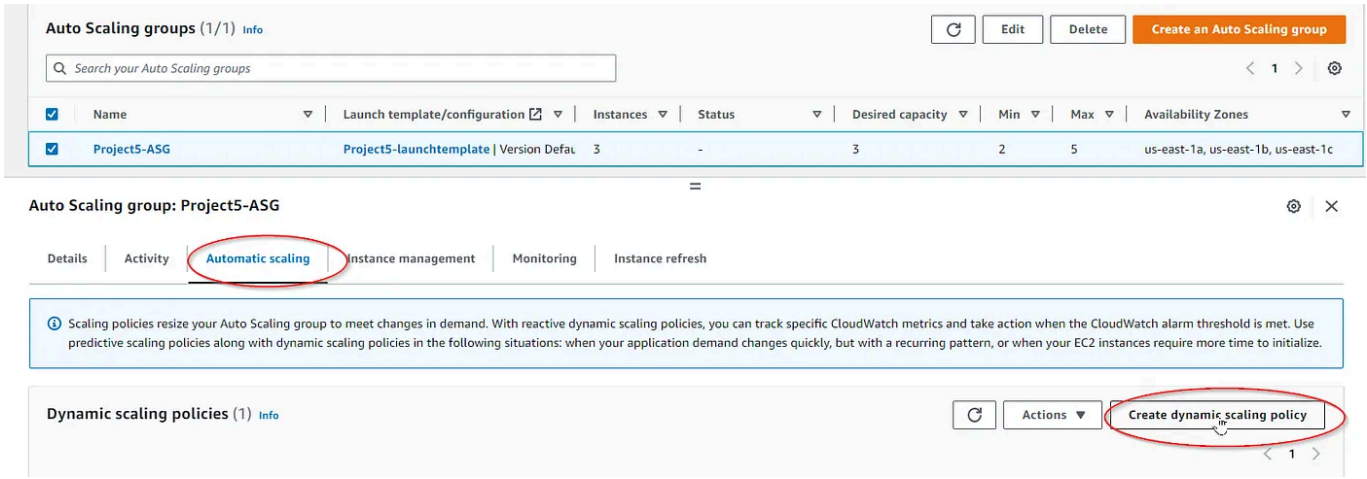

3. Enter 50 for the target value -> Create

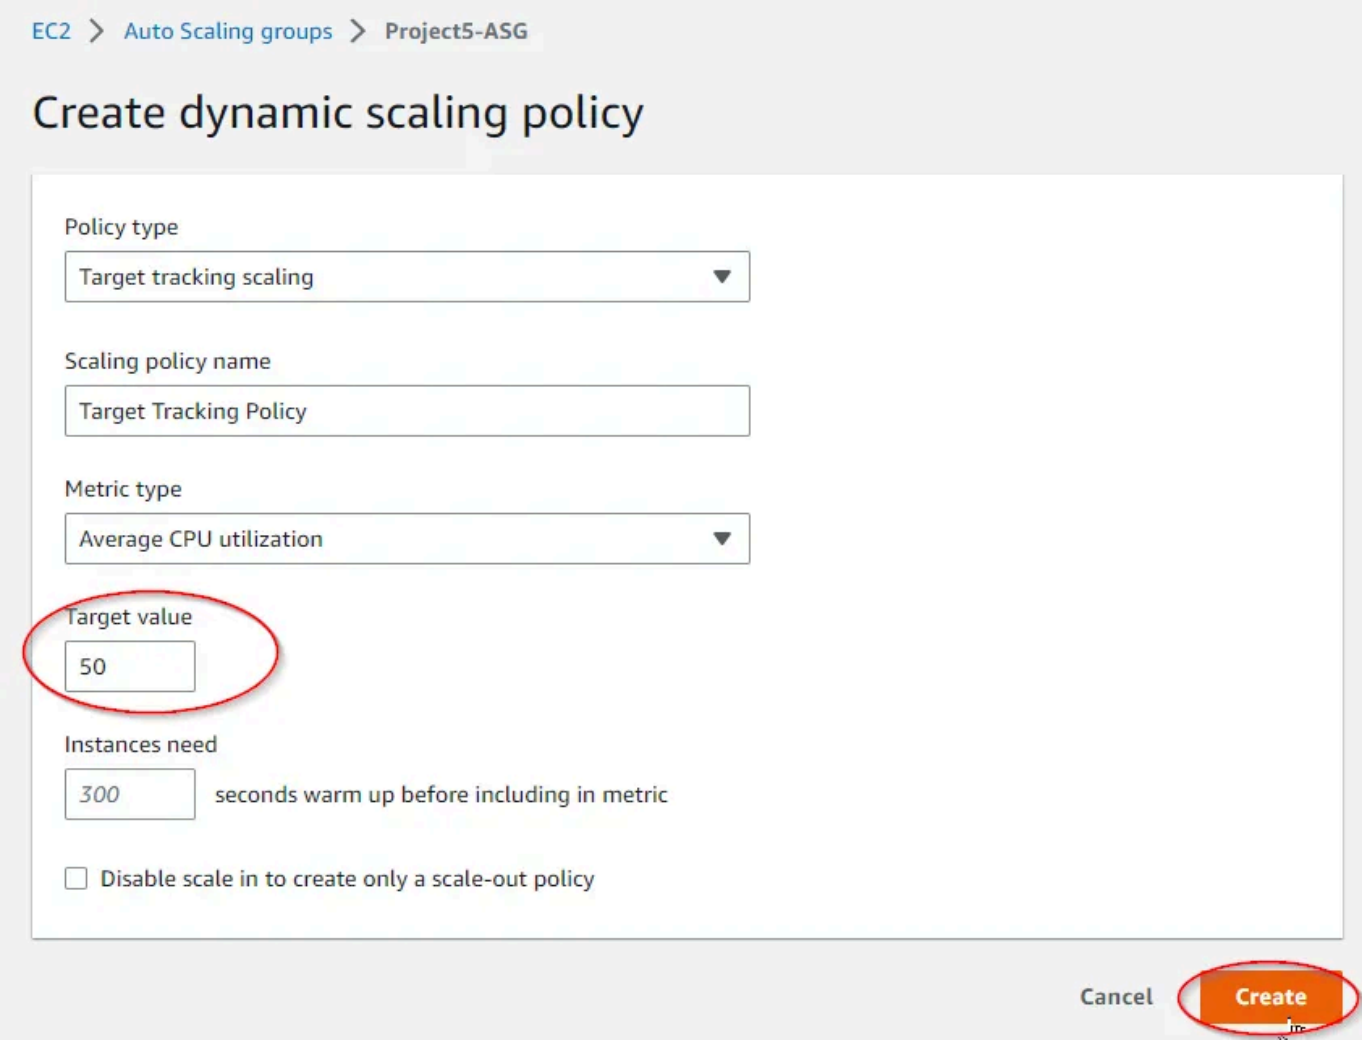

4. SSH into one of your instances and run the following commands to install a stress utility

```
sudo amazon-linux-extras install epel -y
sudo yum install stress -y
```

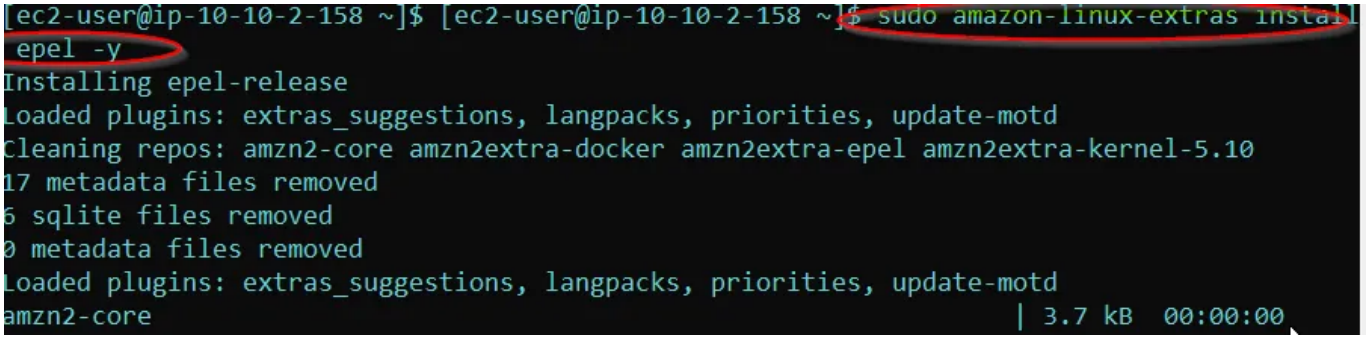

```
[ec2-user@ip-10-10-2-158 ~]$ sudo yum install stress -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
216 packages excluded due to repository priority protections
Resolving Dependencies
--> Running transaction check
---> Package stress.x86_64 0:1.0.4-16.el7 will be installed
--> Finished Dependency Resolution
```

5. Once installed, CPU load can be generated using Stress by running:

```
stress --cpu 1 --timeout 300
```

6. Instance surpassed 60% and generated a new instance

Instance state = running | Clear filters

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone
-	i-0f3f49dcbcc076fe5	Running	t2.micro	2/2 checks passed	No alarms	us-east-1b
	i-0fbe9d16bac360af0	Running	t2.micro	Initializing	No alarms	us-east-1c
-	i-0824158e4c187ac5c	Running	t2.micro	2/2 checks passed	No alarms	us-east-1a

Clean up

1. In the search bar type EC2 -> Detach Load balancer -> Delete Auto Scaling Group -> Delete Launch Template -> Terminate Instance

2. In the search bar type VPC -> Delete VPC

In conclusion, by following these steps to create a VPC with subnets, an autoscaling group, and a load balancer, you have set up a highly available web application infrastructure on AWS. This infrastructure will allow for automatic scaling and distribution of traffic to provide a reliable and responsive user experience. Additionally, the security groups you have set up will ensure that only authorized traffic is allowed to access your resources, providing a secure environment for your web application.